

Practical No: 01

Aim: Implementation Of Prolog

Code:

precious(gold)precious(silver).
precious(pearl).

Output:

```
?- [prac1].  
true.  
  
?- precious(gold).  
true.  
  
?- precious(silver).  
true.  
  
?- precious(platinum).  
false.  
  
?- precious(silve).  
false.
```

Code:

female (mary).
father (surya, ramesh).
likes (surya, food).
likes (surya, car).
likes (surya, bike).
likes (surya, book).
likes (shreya, chocolate).
likes (sudha, food).
likes (surya, coffee).
likes (samudra, chocolate).
likes (samanta, X): likes (X, chocolate).

Output:

```
Welcome to SWI-Prolog (threaded, 64 bits, version 9.2.4)  
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.  
Please run ?- license. for legal details.  
  
For online help and background, visit https://www.swi-prolog.org  
For built-in help, use ?- help(Topic). or ?- apropos(Word).  
  
?- [prac1].  
true.  
  
?- likes(surya,food).  
true ;  
false.  
  
?- likes(shreya,book).  
false.  
  
?- likes(samanta,X).  
X = shreya ,  
  
?- likes(samudra,X).  
X = chocolate.  
  
?-
```

Practical No: 02

Aim: Implementation of Water jug problem using Prolog

Code:

```
water_jug (X, Y): - X>4, Y<3, write ('4L jug overflow'), nl.  
water_jug (X, Y): - X>4, Y>3, write ('3L jug overflow'), nl.  
water_jug (X, Y): - X>4, Y>3, write ('Both jug overflow'), nl.  
water_jug (X, Y): - (X=: =0, Y=: =0, nl, write ('4L:0 & 3L:3 (Action: Fill 3L  
jug.)'), YY is 3,  
water_jug (X, YY));  
(X=: =0, Y=: =0, nl, write ('4L:4 & 3L:0 (Action: Fill 4L jug.)'), XX is 4,  
water_jug (XX, Y));  
(X=: =2, Y=: =0, nl, write ('4L:2 & 3L:0 (Action: Goal Stage reached...)'));  
(X=: =4, Y=: =0, nl, write ('4L:1 & 3L:3 (Action: Pour water from 4L to 3L  
jug.)'), XX is X-3,  
YY is 3, water_jug (XX, YY));  
(X=: =0, Y=: =3, nl, write ('4L:3 & 3L:0 (Action: Pour water from 3L to 4L  
jug.)'), XX is 3, YY  
is 0, water_jug (XX, YY));  
(X=: =1, Y=: =3, nl, write ('4L:1 & 3L:0 (Action: Empty 3L jug.)'), YY is 0,  
water_jug (X, YY));  
(X=: =3, Y=: =0, nl, write ('4L:3 & 3L:3 (Action: Fill 3L jug.)'), YY is 3,  
water_jug (X, YY));  
(X=: =3, Y=: =3, nl, write ('4L:4 & 3L:2 (Action: Pour water from 3L jug to 4L  
jug until 4L jug is full.)'), XX is X+1, YY is Y-1, water_jug (XX, YY));  
(X=: =1, Y=: =0, nl, write ('4L:0 & 3L:1 (Action: Pour water from 4L jug to 3L  
jug.)'), XX is Y,  
YY is X, water_jug (XX, YY));  
(X=: =0, Y=: =1, nl, write ('4L:4 & 3L:1 (Action: Pour water from 4L.)'), XX is 4,  
water_jug (XX, Y));  
(X=: =4, Y=: =1, nl, write ('4L:2 & 3L:3 (Action: Pour water from 4L jug to 3L  
jug until 3L jug  
is full.)'), XX is X-2, YY is Y+2, water_jug (XX, YY));  
(X=: =2, Y=: =3, nl, write ('4L:2 & 3L:0 (Action: Empty 3L jug.)'), YY is 0,  
water_jug (X, YY));  
(X=: = 4, Y=: =2, nl, write ('4L:0 & 3L:2 (Action: Empty 4L jug.)'), XX is 0,  
water_jug (XX, Y));  
(X=: =0, Y=: =2, nl, write ('4L:2 & 3L:0 (Action: Pour water from 3L jug to 4L  
jug.)'), XX is Y,  
YY is X, water_jug (XX, YY)).
```

Output:

```
Welcome to SWI-Prolog (threaded, 64 bits, version 9.2.4)
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.
```

```
For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).
```

```
?- [prac2].
```

```
true.
```

```
?- waterjug(9,0).
```

```
Correct to: "water_jug(9,0)"? yes
```

```
4L jug overflow
```

```
true
```

```
Unknown action: e (h for help)
```

```
Action?
```

```
Unknown action: s (h for help)
```

```
Action? ,
```

```
?- waterjug(0,6).
```

```
Correct to: "water_jug(0,6)"?
```

```
Please answer 'y' or 'n'? yes
```

```
false.
```

```
?- waterjug(6,6).
```

```
Correct to: "water_jug(6,6)"?
```

```
Please answer 'y' or 'n'? yes
```

```
3L jug overflow
```

```
true .
```

```
?- waterjug(0,0).
```

```
Correct to: "water_jug(0,0)"? yes
```

```
4L:0 & 3L:3 (Action : Fill 3L jug.)
```

```
4L:3 & 3L:0 (Action : Pour water from 3L to 4L jug.)
```

```
4L:3 & 3L:3 (Action : Fill 3L jug.)
```

```
4L:4 & 3L:2 (Action : Pour water from 3L jug to 4L jug untill 4L jug  
is full.)
```

```
4L:0 & 3L:2 (Action : Empty 4L jug .)
```

```
4L:2 & 3L:0 (Action : Pour water from 3L jug to 4L jug.)
```

```
4L:2 & 3L:0 (Action : Goal Stage reached...)
```

```
true .
```

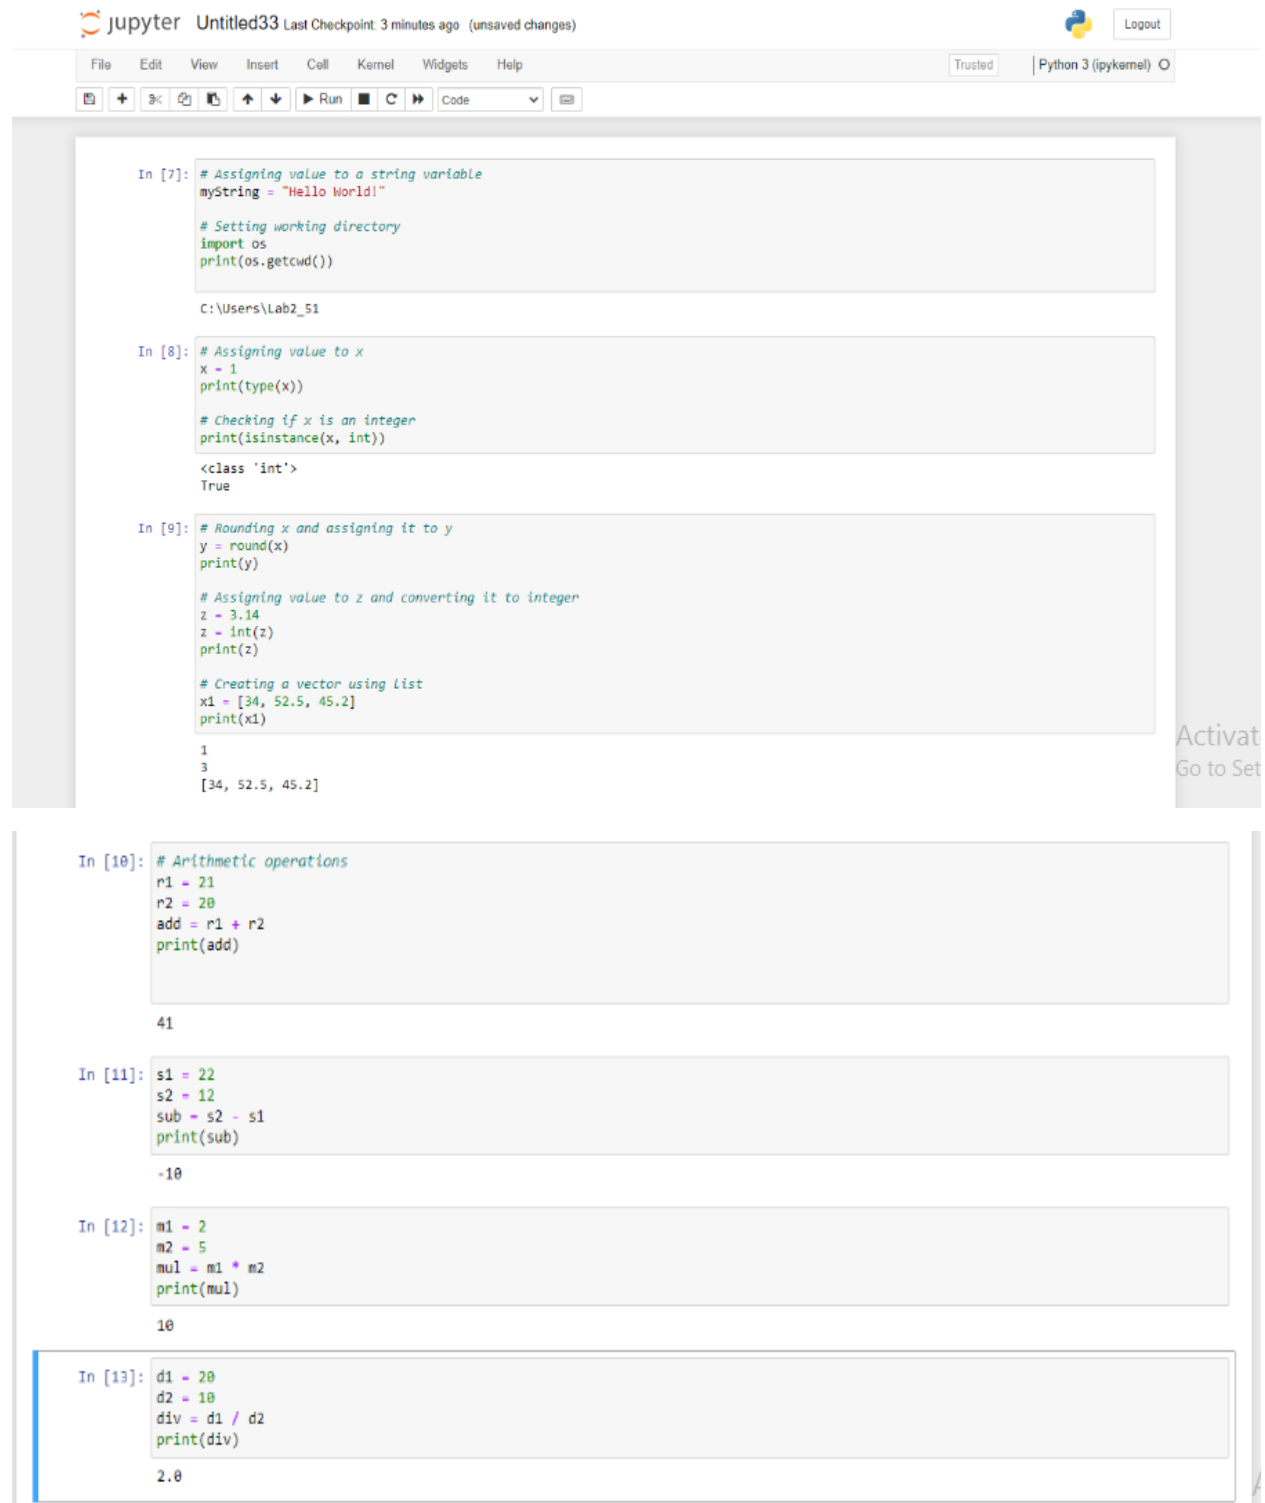
```
?-
```

```
|
```

Practical No:3

Aim: Introduction to Python Programming

Output:



The screenshot displays a Jupyter Notebook interface with the following components:

- Header:** "jupyter Untitled33 Last Checkpoint: 3 minutes ago (unsaved changes)" and a "Logout" button.
- Menu Bar:** File, Edit, View, Insert, Cell, Kernel, Widgets, Help.
- Toolbar:** Includes icons for file operations, cell navigation, and execution, along with a "Code" dropdown menu.
- Code Cells and Outputs:**
 - Cell 7:** Assigns a string to `myString`, imports `os`, and prints the current directory. Output: `C:\Users\Lab2_51`.
 - Cell 8:** Assigns `x = 1`, prints its type, checks if it's an integer, and prints the result. Output: `<class 'int'>` and `True`.
 - Cell 9:** Rounds `x` to `y`, assigns `z = 3.14`, converts it to an integer, and creates a list `x1`. Output: `1`, `3`, and `[34, 52.5, 45.2]`.
 - Cell 10:** Performs arithmetic operations: `r1 = 21`, `r2 = 20`, `add = r1 + r2`. Output: `41`.
 - Cell 11:** Performs arithmetic operations: `s1 = 22`, `s2 = 12`, `sub = s2 - s1`. Output: `-10`.
 - Cell 12:** Performs arithmetic operations: `m1 = 2`, `m2 = 5`, `mul = m1 * m2`. Output: `10`.
 - Cell 13:** Performs arithmetic operations: `d1 = 20`, `d2 = 10`, `div = d1 / d2`. Output: `2.0`.

Practical No:4

Aim: Introduction to Python Libraries

Output:

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df = pd.read_csv('C:\\Users\\USER\\Desktop\\eda\\mtcars.csv')
```

```
In [3]: print(df)
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	

	gear	carb
0	4	4
1	4	4
2	4	1
3	3	1
4	3	2
5	3	1
6	3	4
7	4	2
8	4	2
9	4	4
10	4	4
11	3	3
12	3	3

```
In [5]: df.head()
```

Out[5]:

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1

```
In [ ]: df.head(10)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32 entries, 0 to 31
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   model       32 non-null    object
1   mpg         32 non-null    float64
2   cyl         32 non-null    int64
3   disp        32 non-null    float64
4   hp          32 non-null    int64
5   drat        32 non-null    float64
6   wt          32 non-null    float64
7   qsec        32 non-null    float64
8   vs          32 non-null    int64
9   am          32 non-null    int64
10  gear        32 non-null    int64
11  carb        32 non-null    int64
dtypes: float64(5), int64(6), object(1)
memory usage: 3.1+ KB
```

```
In [9]: df.isnull()
```

Out[9]:

[illegible]

In [10]: `df.isnull().sum()`

Out[10]:

```
model    0
mpg      0
cyl      0
disp     0
hp       0
drat     0
wt       0
qsec     0
vs       0
am       0
gear     0
carb     0
dtype: int64
```

In [11]: `df.tail()`

Out[11]:

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.6	1	1	4	2

In [12]: `df.describe()`

Out[12]:

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750	0.437500	0.406250	3.687500	2.812500
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457	1.786943	0.504016	0.498991	0.737804	1.615548
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000	14.500000	0.000000	0.000000	3.000000	1.000000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250	16.892500	0.000000	0.000000	3.000000	2.000000
50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000	17.710000	0.000000	0.000000	4.000000	2.000000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000	18.900000	1.000000	1.000000	4.000000	4.000000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000	22.900000	1.000000	1.000000	5.000000	8.000000

In [13]: `df.size`

Out[13]: 384

In [14]: `df.shape`

Out[14]: (32, 12)

In [15]: `df.ndim`

Out[15]: 2

In [16]: `df.at[4, 'model']`

Out[16]: 'Hornet Sportabout'

In [17]: `df.at[4, 'disp']`

Out[17]: 360.0

In [18]: `df.iat[3,4]`

Out[18]: 110

In [20]: `df.loc[:, 'model']`

```
Out[20]: 0      Mazda RX4
1      Mazda RX4 Wag
2      Datsun 710
3      Hornet 4 Drive
4      Hornet Sportabout
5      Valiant
6      Duster 360
7      Merc 240D
8      Merc 230
9      Merc 280
10     Merc 280C
11     Merc 450SE
12     Merc 450SL
13     Merc 450SLC
14     Cadillac Fleetwood
15     Lincoln Continental
16     Chrysler Imperial
17     Fiat 128
18     Honda Civic
19     Toyota Corolla
20     Toyota Corona
21     Dodge Challenger
22     AMC Javelin
23     Camaro Z28
24     Pontiac Firebird
25     Fiat X1-9
26     Porsche 914-2
27     Lotus Europa
28     Ford Pantera L
29     Ferrari Dino
30     Maserati Bora
31     Volvo 142E
Name: model, dtype: object
```

```
In [23]: df.iloc[0:5,0:2]
```

```
Out[23]:
```

	model	mpg
0	Mazda RX4	21.0
1	Mazda RX4 Wag	21.0
2	Datsun 710	22.8
3	Hornet 4 Drive	21.4
4	Hornet Sportabout	18.7

```
In [26]: df.iloc[22:32,:]
```

```
Out[26]:
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

```
In [27]: df.dtypes
```

```
Out[27]: model      object
mpg         float64
cyl          int64
disp        float64
hp           int64
drat        float64
wt          float64
qsec        float64
vs           int64
am           int64
gear        int64
carb        int64
dtype: object
```



```
In [28]: df['model'].dtype
```

```
Out[28]: dtype('O')
```

```
In [29]: df.axes
```

```
Out[29]: [RangeIndex(start=0, stop=32, step=1),  
         Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',  
               'gear', 'carb'],  
               dtype='object')]
```

```
In [30]: df.columns
```

```
Out[30]: Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',  
               'gear', 'carb'],  
               dtype='object')
```

```
In [31]: df['hp'].std()
```

```
Out[31]: 68.56286848932059
```

```
In [32]: df['mpg'].mean()
```

```
Out[32]: 20.090625000000003
```

```
In [33]: df['mpg'].median()
```

```
Out[33]: 19.2
```

```
In [34]: df['hp'].describe()
```

```
Out[34]: count      32.000000  
         mean      146.687500  
         std       68.562868  
         min       52.000000  
         25%       96.500000  
         50%      123.000000  
         75%      180.000000  
         max      335.000000  
         Name: hp, dtype: float64
```

Practical No:5

Aim: Program to Implement Linear Regression

Output:

```
[1]: import pandas as pd
```

```
[2]: import numpy as np
```

```
[3]: from sklearn import linear_model
```

```
[4]: import matplotlib.pyplot as plt
```

```
[5]: df = pd.read_csv(r'C:\Users\nisha\OneDrive\Desktop\area.csv')
```


```
[6]: df
```

```
[6]:
```

	area	price
0	2600	550000
1	3000	565000
2	3200	610000
3	3600	680000
4	4000	725000

```
[7]: plt.xlabel('area')
```

```
[7]: Text(0.5, 0, 'area')
```




```
[8]: plt.ylabel('area')
```

```
[8]: Text(0, 0.5, 'area')
```



```
[13]: plt.scatter(df.area,df.price,color='red',marker='+')
```

```
[13]: <matplotlib.collections.PathCollection at 0x19db6be6410>
```



```
[14]: new_df = df.drop('price',axis='columns')
```

```
[15]: new_df
```

```
[15]:
```

	area
0	2600
1	3000
2	3200
3	3600
4	4000

```
[16]: price = df.price
```

```
[17]: price
```

```
[17]:
```

0	550000
1	565000
2	610000
3	680000
4	725000

Name: price, dtype: int64

```
[18]: reg = linear_model.LinearRegression()
```

```
[19]: reg.fit(new_df,price)
```

```
[19]:
```

LinearRegression

LinearRegression()

```
[20]: reg.predict([[3300]])
```

```
C:\Users\nisha\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

```
[20]: array([628715.75342466])
```

```
[21]: reg.coef_
```

```
[21]: array([135.78767123])
```

```
[22]: reg.intercept_
```

```
[22]: 180616.43835616432
```

```
[23]: 3300*135.78767123 + 180616.43835616432
```

```
[23]: 628715.7534151643
```

```
[24]: reg.predict([[5000]])
```

```
C:\Users\nisha\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```


```
[24]: array([859554.79452055])
```

```
[25]: plt.xlabel('area',fontsize=20)
```

```
[24]: array([859554.79452055])
```


```
[25]: plt.xlabel('area',fontsize=20)
```

```
[25]: Text(0.5, 0, 'area')
```




```
[26]: plt.ylabel('price',fontsize=20)
```

```
[26]: Text(0, 0.5, 'price')
```




```
[27]: plt.scatter(df.area,df.price,color='red',marker='+')
```

```
[27]: <matplotlib.collections.PathCollection at 0x19db72ee490>
```




```
[28]: plt.plot(df.area,reg.predict(df[['area']]),color='blue')
```

```
[28]: <matplotlib.lines.Line2D at 0x19db4d44810>
```



```
[29]: plt.xlabel('area',fontsize=20)  
plt.ylabel('price',fontsize=20)  
plt.scatter(df.area,df.price,color='red',marker='+')  
plt.plot(df.area,reg.predict(df[['area']]),color='blue')
```

```
[29]: <matplotlib.lines.Line2D at 0x19db7589890>
```



```
[ ]:
```

Practical No:6

Aim: Program to Implement Logistic Regression

Output:

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: credit_df = pd.read_csv('C:/Users/nisha/OneDrive/Desktop/CreditRisk.csv')
```

```
[3]: credit_df.shape
```

```
[3]: (614, 13)
```

```
In [4]: credit_df.head()
```

Out[4]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [5]: credit_df.tail()
```

Out[5]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71	360.0	1.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40	180.0	1.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253	360.0	1.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187	360.0	1.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133	360.0	0.0	

```
In [6]: credit_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Loan_ID             614 non-null   object 
 1   Gender              601 non-null   object 
 2   Married             611 non-null   object 
 3   Dependents          599 non-null   object 
 4   Education            614 non-null   object 
 5   Self_Employed       582 non-null   object 
 6   ApplicantIncome     614 non-null   int64  
 7   CoapplicantIncome   614 non-null   float64 
 8   LoanAmount          614 non-null   int64  
 9   Loan_Amount_Term    600 non-null   float64 
10   Credit_History       564 non-null   float64 
11   Property_Area       614 non-null   object 
12   Loan_Status         614 non-null   int64  
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

In [7]: `credit_df.describe()`

Out[7]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.000000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.000000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.000000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.000000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.000000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.000000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000	1.000000

In [8]: `credit_df.Loan_Status.value_counts()`


Out[8]: 1 422
0 192
Name: Loan_Status, dtype: int64

In [9]: `credit_df.groupby(['Education', 'Loan_Status']).Education.count()`

Out[9]: Education Loan_Status
Graduate 0 140
1 340
Not Graduate 0 52
1 82
Name: Education, dtype: int64

In [10]: `sns.barplot(y = 'Credit_History', x='Loan_Status', hue='Education', data=credit_df)`

Out[10]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>



```
In [11]: 100 * credit_df.isnull().sum() / credit_df.shape[0]
```

```
Out[11]: Loan_ID          0.000000
Gender          2.117264
Married         0.488599
Dependents      2.442997
Education       0.000000
Self_Employed   5.211726
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount      0.000000
Loan_Amount_Term 2.280130
Credit_History  8.143322
Property_Area   0.000000
Loan_Status     0.000000
dtype: float64
```

```
In [12]: DF=credit_df.drop(credit_df.columns[0],axis=1)
```

```
In [13]: DF.head()
```

```
Out[13]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

```
In [14]: object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
```

```
In [15]: for column in object_columns:
majority = DF[column].value_counts().iloc[0]
DF[column].fillna(majority, inplace=True)
```

```
In [17]: for column in numeric_columns:
mean = DF[column].mean()
DF[column].fillna(mean, inplace=True)
```

```
In [18]: DF.head()
```

```
Out[18]:
```

	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	1
	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	0
	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	1
	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	1
	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	1


```
In [20]: DF[object_columns].Property_Area #Categorical Columns
```

```
Out[20]: 0      Urban
          1      Rural
          2      Urban
          3      Urban
          4      Urban
          ...
          609    Rural
          610    Rural
          611    Urban
          612    Urban
          613    Semiurban
          Name: Property_Area, Length: 614, dtype: object
```

```
In [21]: DF[object_columns].Property_Area.head()
```

```
Out[21]: 0      Urban
          1      Rural
          2      Urban
          3      Urban
          4      Urban
          Name: Property_Area, dtype: object
```

```
In [22]: Df_dummy = pd.get_dummies(DF, columns=object_columns)
```

```
In [23]: Df_dummy.head()
```

```
Out[23]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_398
0	5849	0.0	0	360.0	1.0	1	0	0	1	0
1	4583	1508.0	128	360.0	1.0	0	0	0	1	0
2	3000	0.0	66	360.0	1.0	1	0	0	1	0
3	2583	2358.0	120	360.0	1.0	1	0	0	1	0
4	6000	0.0	141	360.0	1.0	1	0	0	1	0

5 rows x 25 columns

```
In [24]: Df_dummy.shape
```

```
Out[24]: (614, 25)
```

```
In [25]: from sklearn.model_selection import train_test_split
          from sklearn.linear_model import LogisticRegression
          from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [29]: x = Df_dummy.drop('Loan_Status', axis=1)
          y = Df_dummy.Loan_Status
          train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=42)
```

```
In [30]: train_x.shape, test_x.shape
```

```
Out[30]: ((429, 24), (185, 24))
```

```
model = LogisticRegression() #LogisticRegression_Model
```

```
model.fit(train_x, train_y)
```

```
train_y_hat = model.predict(train_x)  
test_y_hat = model.predict(test_x)
```

```
print('train_accuracy', accuracy_score(train_y, train_y_hat))  
print('test accuracy', accuracy_score(test_y, test_y_hat))
```

```
train_accuracy 0.8205128205128205  
test accuracy 0.7837837837837838
```

```
In [35]: print(confusion_matrix(train_y, train_y_hat))
```

```
[[ 57  70]  
 [  7 295]]
```

```
In [36]: print(confusion_matrix(test_y, test_y_hat))
```

```
[[ 27  38]  
 [  2 118]]
```

```
In [37]: test_y.value_counts()
```

```
Out[37]: 1    120  
         0     65  
         Name: Loan_Status, dtype: int64
```

```
In [38]: pd.Series(test_y_hat).value_counts()
```

```
Out[38]: 1    156  
         0     29  
         dtype: int64
```

```
In [39]: (57 + 295) / train_y.shape[0]
```


```
Out[39]: 0.8205128205128205
```

```
In [40]: print(classification_report(test_y, test_y_hat))
```

	precision	recall	f1-score	support
0	0.93	0.42	0.57	65
1	0.76	0.98	0.86	120
accuracy			0.78	185
macro avg	0.84	0.70	0.71	185
weighted avg	0.82	0.78	0.76	185

```
In [41]: x = np.linspace(-10, 10, 100)
y = 1 / (1 + np.exp(-x)) #Sigmoid
plt.plot(x, y)
```

Out[41]: [<matplotlib.lines.Line2D at 0x1d457f4c460>]



Practical No:7

Aim: Implementation Of KNN Classification.

Jupyter KNN Last Checkpoint: 3 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[1]: import pandas as pd
from sklearn.datasets import load_iris
iris=load_iris()

[2]: iris.feature_names

[2]: ['sepal length (cm)',
'sepal width (cm)',
'petal length (cm)',
'petal width (cm)']

[3]: iris.target_names

[3]: array(['setosa', 'versicolor', 'virginica'], dtype='<U10')

[5]: df=pd.DataFrame(iris.data,columns=iris.feature_names)
df.head()

[5]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```
[6]: df['target']=iris.target
df.head()

[6]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

Jupyter KNN Last Checkpoint: 4 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[6]: df['target']=iris.target
df.head()

[6]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
[7]: df[df.target==1].head()

[7]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
50	7.0	3.2	4.7	1.4	1
51	6.4	3.2	4.5	1.5	1
52	6.9	3.1	4.9	1.5	1
53	5.5	2.3	4.0	1.3	1
54	6.5	2.8	4.6	1.5	1

```
[8]: df['flower_name']=df.target.apply(lambda x:iris.target_names[x])
df.head()

[8]:
```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa

```
[9]: df0=df[:50]
    df1=df[50:100]
    df2=df[100:]

[10]: import matplotlib.pyplot as plt
    %matplotlib inline

[12]: plt.xlabel('Sepal Length')
    plt.ylabel('Sepal Width')
    plt.scatter(df0['sepal length (cm)'],df0['sepal width (cm)'],color="green",marker='+')
    plt.scatter(df1['sepal length (cm)'],df1['sepal width (cm)'],color="blue",marker='.')


[12]: <matplotlib.collections.PathCollection at 0x15d77525310>
```



```
[13]: plt.xlabel('Petal Length')
    plt.ylabel('Petal Width')
    plt.scatter(df0['petal length (cm)'],df0['petal width (cm)'],color="green",marker='+')
```

```
plt.scatter(df0['petal length (cm)'],df0['petal width (cm)'],color="green",marker='+')
plt.scatter(df1['petal length (cm)'],df1['petal width (cm)'],color="blue",marker='.')
plt.scatter(df2['petal length (cm)'],df2['petal width (cm)'],color="red",marker='*')
```

```
[13]: <matplotlib.collections.PathCollection at 0x15d7caccf50>
```



```
[14]: from sklearn.model_selection import train_test_split
```

```
[15]: x=df.drop(['target','flower_name'],axis='columns')
    y=df.target
    x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=1)
```

```
[16]: len(x_train)
```

```
[16]: 120
```

```
[17]: len(x_test)
```

```
[17]: 20
```

jupyter KNN Last Checkpoint: 6 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[17]: 30

[19]: from sklearn.neighbors import KNeighborsClassifier
      knn=KNeighborsClassifier(n_neighbors=10)

[20]: knn.fit(x_train,y_train)

[20]: KNeighborsClassifier
      KNeighborsClassifier(n_neighbors=10)

[21]: knn.score(x_test,y_test)

[21]: 0.9666666666666667

[22]: knn.predict([[4.8,3.0,1.5,0.3]])

C:\Users\nisha\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names
  warnings.warn(

[22]: array([0])

[23]: from sklearn.metrics import confusion_matrix
      y_pred=knn.predict(x_test)
      cm=confusion_matrix(y_test,y_pred)
      cm

[23]: array([[11,  0,  0],
       [ 0, 12,  1],
       [ 0,  0,  6]], dtype=int64)

[25]: %matplotlib inline
      import matplotlib.pyplot as plt
      import seaborn as sn
      plt.figure(figsize=(7,5))
      sn.heatmap(cm,annot=True)
      plt.xlabel('Predicted')
      plt.ylabel('Truth')

[25]: Text(58.22222222222214, 0.5, 'Truth')
```



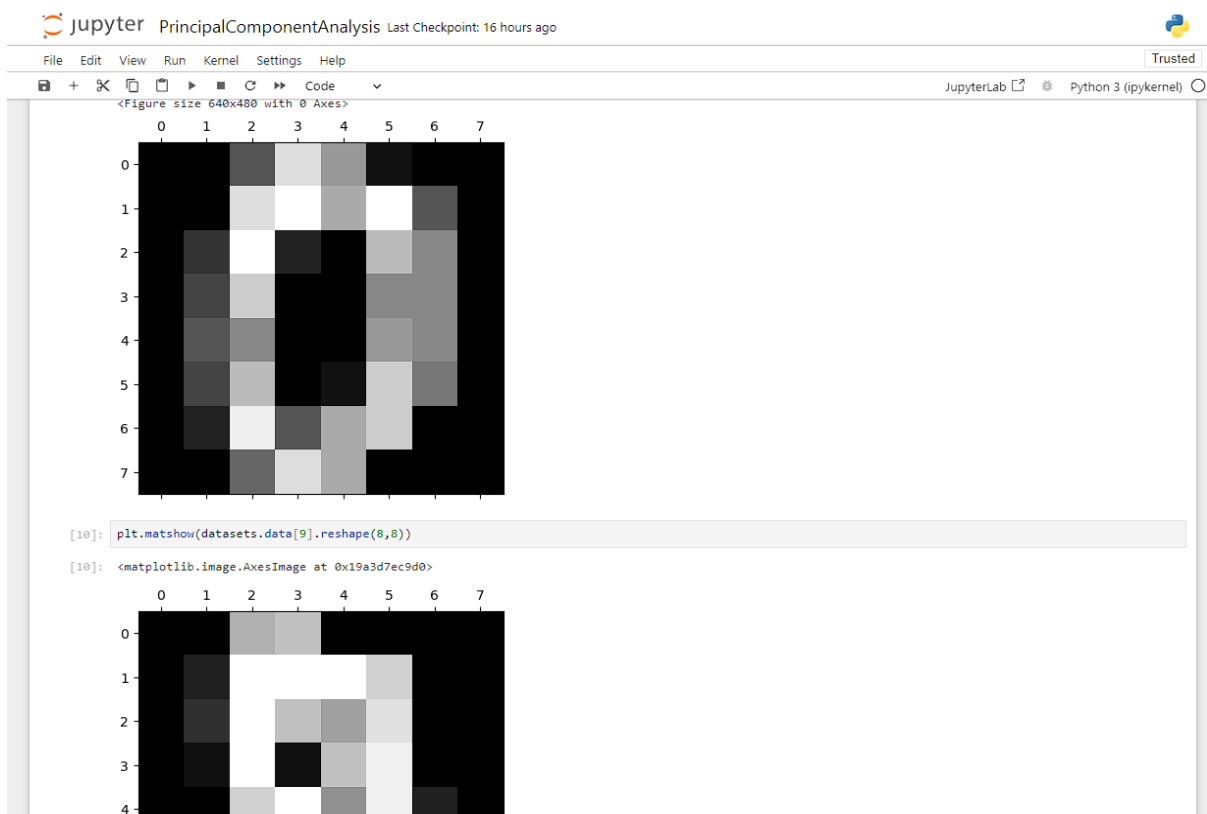
Practical No:8

Aim: Program to Implement Principal Component Analysis

Output:

```
Jupyter PrincipalComponentAnalysis Last Checkpoint: 16 hours ago
File Edit View Run Kernel Settings Help
+ - X Copy Paste Run Code
JupyterLab Python 3 (ipykernel)

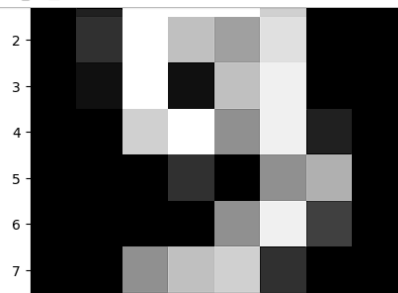
[1]: from sklearn.datasets import load_digits
[2]: import pandas as pd
[3]: datasets=load_digits()
[4]: datasets.keys()
[4]: dict_keys(['data', 'target', 'frame', 'feature_names', 'target_names', 'images', 'DESCR'])
[5]: datasets.data.shape
[5]: (1797, 64)
[6]: datasets.data[0]
[6]: array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
          15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
          12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
           0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
          10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
[7]: datasets.data[0].reshape(8,8)
[7]: array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
          [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
          [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
          [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
          [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
          [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
          [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
          [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
[9]: from matplotlib import pyplot as plt
      %matplotlib inline
      plt.gray()
      plt.matshow(datasets.data[0].reshape(8,8))
[9]: <matplotlib.image.AxesImage at 0x19a3c5b46d0>
      <Figure size 640x480 with 0 Axes>
```



Jupyter PrincipalComponentAnalysis Last Checkpoint: 16 hours ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)



```
[11]: datasets.target[:5]
```

```
[11]: array([0, 1, 2, 3, 4])
```

```
[13]: df=pd.DataFrame(datasets.data,columns=datasets.feature_names)
df.head()
```

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_6	pixel_6_7	pixel_7_0	pixel_7_1	pixel_7_2	pixel_7_3
0	0.0	0.0	5.0	13.0	9.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	6.0
1	0.0	0.0	0.0	12.0	13.0	5.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	4.0	15.0	12.0	0.0	0.0	0.0	0.0	...	5.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	7.0	15.0	13.0	1.0	0.0	0.0	0.0	8.0	...	9.0	0.0	0.0	0.0	0.0	7.0
4	0.0	0.0	0.0	1.0	11.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 64 columns

```
[14]: datasets.target
```

```
[14]: array([0, 1, 2, ..., 8, 9, 8])
```

Jupyter PrincipalComponentAnalysis Last Checkpoint: 16 hours ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[15]: df.describe()
```

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_6	pixel_6_7
count	1797.0	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	...	1797.000000	1797.000000
mean	0.0	0.303840	5.204786	11.835838	11.848080	5.781859	1.362270	0.129661	0.005565	1.993879	...	3.725097	0.206455
std	0.0	0.907192	4.754826	4.248842	4.287388	5.666418	3.325775	1.037383	0.094222	3.196160	...	4.919406	0.984401
min	0.0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000
25%	0.0	0.000000	1.000000	10.000000	10.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000
50%	0.0	0.000000	4.000000	13.000000	13.000000	4.000000	0.000000	0.000000	0.000000	0.000000	...	1.000000	0.000000
75%	0.0	0.000000	9.000000	15.000000	15.000000	11.000000	0.000000	0.000000	0.000000	3.000000	...	7.000000	0.000000
max	0.0	8.000000	16.000000	16.000000	16.000000	16.000000	16.000000	15.000000	2.000000	16.000000	...	16.000000	13.000000

8 rows × 64 columns

```
[16]: x=df
y=datasets.target
```

```
[17]: from sklearn.preprocessing import StandardScaler
```

```
[19]: scaler=StandardScaler()
x_scaled=scaler.fit_transform(x)
x_scaled
```

```
[19]: array([[ 0.          , -0.33501649, -0.04308102, ..., -1.14664746,
          -0.5056698 , -0.19600752],
        [ 0.          , -0.33501649, -1.09493684, ...,  0.54856067,
          -0.5056698 , -0.19600752],
        [ 0.          , -0.33501649, -1.09493684, ...,  1.56568555,
          1.6951369 , -0.19600752],
        ...,
        [ 0.          , -0.33501649, -0.88456568, ..., -0.12952258,
          -0.5056698 , -0.19600752],
        [ 0.          , -0.33501649, -0.67419451, ...,  0.8876023 ,
          -0.5056698 , -0.19600752],
```


Jupyter PrincipalComponentAnalysis Last Checkpoint: 16 hours ago

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

```

-0.5056698, -0.19600752],
[ 0.        , -0.33501649, 1.00877481, ..., 0.8876023 ,
 -0.26113572, -0.19600752]])

[20]: from sklearn.model_selection import train_test_split

[21]: x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2, random_state=30)

[22]: from sklearn.linear_model import LogisticRegression

[23]: model = LogisticRegression()
model.fit(x_train, y_train)
model.score(x_test, y_test)

[23]: 0.9722222222222222

[24]: from sklearn.decomposition import PCA

[26]: pca = PCA(0.95)
x_pca = pca.fit_transform(x)
x_pca.shape

[26]: (1797, 29)

[27]: pca.explained_variance_ratio_

[27]: array([0.14890594, 0.13618771, 0.11794594, 0.08409979, 0.05782415,
 0.0491691 , 0.04315987, 0.03661373, 0.03353248, 0.03078806,
 0.02372341, 0.02272697, 0.01821863, 0.01773855, 0.01467181,
 0.01409716, 0.01318589, 0.01248138, 0.01017718, 0.00905617,
 0.00889538, 0.00797123, 0.00767493, 0.00722904, 0.00695889,
 0.00596081, 0.00575615, 0.00515158, 0.0048954 ])

[28]: pca.n_components_

[28]: 29

[29]: x_train_pca, x_test_pca, y_train, y_test = train_test_split(x_pca, y, test_size=0.2, random_state=30)

[30]: from sklearn.linear_model import LogisticRegression

[31]: model = LogisticRegression(max_iter=1000)

```

Jupyter PrincipalComponentAnalysis Last Checkpoint: 16 hours ago

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

```

[29]: x_train_pca, x_test_pca, y_train, y_test = train_test_split(x_pca, y, test_size=0.2, random_state=30)

[30]: from sklearn.linear_model import LogisticRegression

[31]: model = LogisticRegression(max_iter=1000)

[36]: model = LogisticRegression(max_iter=1000)
model.fit(x_train_pca, y_train)
model.score(x_test_pca, y_test)

[36]: 0.6083333333333333

[37]: pca = PCA(n_components=2)
x_pca = pca.fit_transform(x)
x_pca.shape

[37]: (1797, 2)

[38]: pca.explained_variance_ratio_

[38]: array([0.14890594, 0.13618771])

[39]: x_train_pca, x_test_pca, y_train, y_test = train_test_split(x_pca, y, test_size=0.2, random_state=30)

[40]: model = LogisticRegression(max_iter=1000)
model.fit(x_train_pca, y_train)
model.score(x_test_pca, y_test)

[40]: 0.6083333333333333

[ ]:

```

Practical No: 9

Aim: Program to Implement K-Means Algorithm

Output:

```
In [3]: import pandas as pd
        from sklearn.cluster import KMeans
        from sklearn.preprocessing import MinMaxScaler
        import matplotlib.pyplot as plt
        %matplotlib inline
```


```
In [5]: df=pd.read_csv('C:/Users/LAB2_41/Documents/Siddharth_63/income.csv')
        df.head()
```

```
Out[5]:
```

	Name	Age	Income(\$)
0	Rob	27	70000
1	Michael	29	90000
2	Mohan	29	61000
3	Ismail	28	60000
4	Kory	42	150000

```
In [6]: plt.scatter(df.Age,df['Income($)'])
        plt.xlabel('Age')
        plt.ylabel('Income($)')
```

```
Out[6]: Text(0, 0.5, 'Income($)')
```



```
In [7]: km=KMeans(n_clusters=3)
        y_predicted=km.fit_predict(df[['Age', 'Income($)']])
        y_predicted
```

```
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
  warnings.warn(
```

```
Out[7]: array([2, 2, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 2, 2, 0])
```

```
In [8]: df['clusters']=y_predicted
        df.head()
```

```
Out[8]:
```

	Name	Age	Income(\$)	clusters
0	Rob	27	70000	2
1	Michael	29	90000	2
2	Mohan	29	61000	0
3	Ismail	28	60000	0
4	Kory	42	150000	1


In [9]: km.cluster_centers_

Out[9]: array([[3.29090909e+01, 5.61363636e+04],
[3.82857143e+01, 1.50000000e+05],
[3.40000000e+01, 8.05000000e+04]])

In [13]: df1=df[df.clusters==0]
df2=df[df.clusters==1]
df3=df[df.clusters==2]
plt.scatter(df1.Age,df1['Income(\$)',color='green'])
plt.scatter(df2.Age,df2['Income(\$)',color='red'])
plt.scatter(df3.Age,df3['Income(\$)',color='black'])
plt.xlabel('Age')
plt.ylabel('Income(\$)')
plt.legend()

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend () is called with no argument.

Out[13]: <matplotlib.legend.Legend at 0x2011fb12da0>



Activate Win
Go to Settings

In [14]: scaler=MinMaxScaler()

scaler.fit(df[['Income(\$)']])
df['Income(\$)']=scaler.transform(df[['Income(\$)']])

scaler.fit(df[['Age']])
df['Age']=scaler.transform(df[['Age']])

In [16]: df.head(10)


Out[16]:

	Name	Age	Income(\$)	clusters
0	Rob	0.058824	0.213675	2
1	Michael	0.176471	0.384615	2
2	Mohan	0.176471	0.136752	0
3	Ismail	0.117647	0.128205	0
4	Kory	0.941176	0.897436	1
5	Gautam	0.764706	0.940171	1
6	David	0.882353	0.982906	1
7	Andrea	0.705882	1.000000	1
8	Brad	0.588235	0.948718	1
9	Angelina	0.529412	0.726496	1

In [19]: plt.scatter(df.Age,df['Income(\$)'])

Activate Win
Go to Settings

Out[19]: <matplotlib.collections.PathCollection at 0x20123121c90>



Activate Win
Go to Settings tc

```
In [24]: km=KMeans(n_clusters=3)
y_predicted=km.fit_predict(df[['Age','Income($)']])
y_predicted
```

C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
warnings.warn(

Out[24]: array([1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2])

```
In [25]: df['cluster']=y_predicted
df.head()
```

Out[25]:

	Name	Age	Income(\$)	clusters	cluster
0	Rob	0.058824	0.213675	2	1
1	Michael	0.176471	0.384615	2	1
2	Mohan	0.176471	0.136752	0	1
3	Ismail	0.117647	0.128205	0	1
4	Kory	0.941176	0.897436	1	0


```
In [26]: km.cluster_centers_
```

Out[26]: array([[0.72268908, 0.8974359],
[0.1372549 , 0.11633428],
[0.85294118, 0.2022792]])

Activate Win

```
In [29]: df1=df[df.clusters==0]
df2=df[df.clusters==1]
df3=df[df.clusters==2]
plt.scatter(df1.Age,df1['Income($)'],color='green')
plt.scatter(df2.Age,df2['Income($)'],color='red')
plt.scatter(df3.Age,df3['Income($)'],color='black')
plt.scatter(km.cluster_centers_[0,0],km.cluster_centers_[0,1],color='purple',marker='*',label='centroid')
plt.legend()
```

Out[29]: <matplotlib.legend.Legend at 0x201232ab910>




Activate Wi
Go to Settings 1

```
In [30]: sse=[]
          k_rng=range(1,10)
          for k in k_rng:
              km=KMeans(n_clusters=k)
              km.fit(df[['Age', 'Income($)']])
              sse.append(km.inertia_)
```

```
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.  
warnings.warn(  
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.  
warnings.warn(  
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.  
warnings.warn(  
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.  
warnings.warn(  
C:\Users\LAB2_41\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1334: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable
```

```
In [32]: plt.xlabel('K')
plt.ylabel('Sum of squared error')
plt.plot(k_rng,sse)
plt.title('Elbow Method')
```

```
Out[32]: Text(0.5, 1.0, 'Elbow Method')
```



Activate Wi-Fi
Go to Settings 1

Practical No: 10

Aim: Implementation Of Support Vector Machine.

Output:

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: credit_df=pd.read_csv(r'C:\Users\nisha\OneDrive\Desktop\SEM 2\AI ML LAB\CreditRisk.csv')
credit_df.head()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
[4]: credit_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column              Non-Null Count  Dtype  
---  --
 0   Loan_ID             614 non-null    object  
 1   Gender              601 non-null    object  
 2   Married             611 non-null    object  
 3   Dependents          599 non-null    object  
 4   Education           614 non-null    object  
 5   Self_Employed       582 non-null    object  
 6   ApplicantIncome     614 non-null    int64   
 7   CoapplicantIncome   614 non-null    float64  
 8   LoanAmount          614 non-null    int64   
 9   Loan_Amount_Term    600 non-null    float64  
10   Credit_History       564 non-null    float64  
11   Property_Area       614 non-null    object  
12   Loan_Status         614 non-null    int64   
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

```
[5]: credit_df.Loan_Status.value_counts()

Loan_Status
1      422
0      192
Name: count, dtype: int64
```

```
[6]: credit_df.describe()
```


	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.00000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.00000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.00000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.00000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.00000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.00000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000	1.000000

```
[7]: credit_df.groupby(['Education','Loan_Status']).Education.count()

Education    Loan_Status
Graduate     0             140
              1             340
Not Graduate  0             52
              1             82
Name: Education, dtype: int64
```

```
[8]: sns.barplot(y='Credit_History', x='Loan_Status', hue='Education',data = credit_df)
```

```
[8]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```



```
[9]: 100 * credit_df.isnull().sum() / credit_df.shape[0]
```

```
[9]: Loan_ID      0.000000
Gender        2.117264
Married       0.488599
Dependents    2.442997
Education     0.000000
Self_Employed 5.211726
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount    0.000000
Loan_Amount_Term 2.280130
Credit_History 8.143322
Property_Area 0.000000
Loan_Status   0.000000
dtype: float64
```

```
[10]: DF = credit_df.drop('Loan_ID', axis = 1)
```

```
[11]: object_columns = DF.select_dtypes(include = ['object']).columns
numeric_columns = DF.select_dtypes(exclude = ['object']).columns
```

```
[12]: for column in object_columns:
    majority = DF[column].value_counts().iloc[0]
    DF[column].fillna(majority, inplace=True)
```

```
[13]: for column in numeric_columns:
    mean = DF[column].mean()
    DF[column].fillna(mean, inplace=True)
```

```
[14]: mean
```

```
[14]: 0.6872964169381107
```

```
[15]: DF.head()
```

```
[15]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	0
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	1
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	0
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	1
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	0

```
[16]: DF[object_columns].Married
```

```
[16]: 0    No
1    Yes
2    Yes
3    Yes
4    No
...
609  No
610  Yes
611  Yes
612  Yes
613  No
Name: Married, Length: 614, dtype: object
```

```
[17]: DF[object_columns].Property_Area.head()
```

```
[17]: DF[object_columns].Property_Area.head()

[17]: 0    Urban
      1    Rural
      2    Urban
      3    Urban
      4    Urban
      Name: Property_Area, dtype: object

[18]: DF_dummy = pd.get_dummies(DF, columns = object_columns)

[19]: DF_dummy.shape
      (614, 25)

[20]: DF_dummy.head()

[20]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_398	...	De
0	5849	0.0	0	360.0	1.0	1	False	False	True	False	...	
1	4583	1508.0	128	360.0	1.0	0	False	False	True	False	...	
2	3000	0.0	66	360.0	1.0	1	False	False	True	False	...	
3	2583	2358.0	120	360.0	1.0	1	False	False	True	False	...	
4	6000	0.0	141	360.0	1.0	1	False	False	True	False	...	

```

5 rows x 25 columns

[21]: from sklearn.model_selection import train_test_split as TTS
      from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

[22]: X = DF_dummy.drop('Loan_Status', axis=1)
      Y = DF_dummy.Loan_Status
      train_x, test_x, train_y, test_y = TTS(X, Y, test_size = 0.3, random_state=42)

[23]: train_x.shape, test_x.shape

[23]: ((429, 24), (185, 24))

[24]: svm_model = SVC(kernel='rbf', gamma=0.00001, C=1000)

[25]: svm_model.fit(train_x, train_y)

[25]: SVC(C=1000, gamma=1e-05)

[26]: train_y_hat = svm_model.predict(train_x)
      test_y_hat = svm_model.predict(test_x)

[27]: print('-'*20, 'Train', '-'*20)
      print(classification_report(train_y, train_y_hat))
      print('-'*20, 'Test', '-'*20)
      print(classification_report(test_y, test_y_hat))

----- Train -----
              precision    recall  f1-score   support

      0       0.95       0.95       0.95        127
      1       0.98       0.98       0.98        302

   accuracy       0.96
  macro avg       0.96
 weighted avg       0.97

----- Test -----
              precision    recall  f1-score   support

      0       0.36       0.18       0.24         65
      1       0.65       0.82       0.73        120

   accuracy       0.51
  macro avg       0.50
 weighted avg       0.55

[28]: confusion_matrix(train_y, train_y_hat)

[28]: array([[121,  6],
          [ 7, 295]], dtype=int64)

[ ]:
```


Practical No: 11

Aim: Program to Implement Decision Tree

Output:

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: credit_df=pd.read_csv(r"C:\Users\USER\Documents\2155a1ml\CreditRisk.csv")
credit_df.head()
```

```
Out[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0

```
In [5]: credit_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Loan_ID               614 non-null   object  
1   Gender                601 non-null   object  
2   Married               611 non-null   object  
3   Dependents            599 non-null   object  
4   Education              614 non-null   object  
5   Self_Employed         582 non-null   object  
6   ApplicantIncome       614 non-null   int64   
7   CoapplicantIncome     614 non-null   float64  
8   LoanAmount            614 non-null   int64   
9   Loan_Amount_Term      600 non-null   float64  
10  Credit_History         564 non-null   float64  
11  Property_Area         614 non-null   object  
12  Loan_Status           614 non-null   int64   
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

```
In [7]: credit_df.Loan_Status.value_counts()
```

```
Out[7]: Loan_Status
1      422
0      192
Name: count, dtype: int64
```

```
In [8]: credit_df.describe()
```

```
Out[8]:
```


	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.000000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.000000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.000000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.000000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.000000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.000000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000	1.000000

```
In [10]: credit_df.groupby(['Education','Loan_Status']).Education.count()
```

```
Out[10]: Education      Loan_Status
Graduate      0          140
              1          340
Not Graduate   0          52
              1          82
Name: Education, dtype: int64
```

```
In [11]: sns.barplot(y='Credit_History',x='Loan_Status',hue='Education',data=credit_df)
```

```
Out[11]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```



```
In [12]: 100*credit_df.isnull().sum()/credit_df.shape[0]
```

```
Out[12]: Loan_ID          0.000000
Gender          2.117264
Married         0.488599
Dependents      2.442997
Education       0.000000
Self_Employed   5.211726
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount      0.000000
Loan_Amount_Term 2.280130
Credit_History  8.143322
Property_Area   0.000000
Loan_Status     0.000000
dtype: float64
```

```
In [13]: DF = credit_df.drop('Loan_ID', axis = 1)
object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
for column in object_columns:
    majority = DF[column].value_counts().iloc[0]
    DF[column].fillna(majority, inplace=True)
for column in numeric_columns:
    mean = DF[column].mean()
    DF[column].fillna(mean, inplace=True)
```

```
In [14]: mean
```

In [15]: DF.head()

Out[15]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

In [16]: DF[object_columns].Married

Out[16]:

```
0      No
1      Yes
2      Yes
3      Yes
4      No
...
609    No
610    Yes
611    Yes
612    Yes
613    No
Name: Married, Length: 614, dtype: object
```

In [17]: DF[object_columns].Property_Area.head()

Out[17]:

```
0      Urban
1      Rural
2      Urban
3      Urban
4      Urban
Name: Property_Area, dtype: object
```

In [18]: DF_dummy = pd.get_dummies(DF, columns = object_columns)
DF_dummy.shape

Out[18]: (614, 25)

In [19]: DF_dummy.head()

Out[19]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_398
0	5849	0.0	0	360.0	1.0	1	False	False	True	False
1	4583	1508.0	128	360.0	1.0	0	False	False	True	False
2	3000	0.0	66	360.0	1.0	1	False	False	True	False
3	2583	2358.0	120	360.0	1.0	1	False	False	True	False
4	6000	0.0	141	360.0	1.0	1	False	False	True	False

5 rows x 25 columns

In [20]: from sklearn.model_selection import train_test_split as TTS
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

In [24]: X = DF_dummy.drop('Loan_Status',axis=1)
Y = DF_dummy.Loan_Status
train_x, test_x, train_y, test_y=TTS(X,Y,test_size=0.3,random_state=42)

In [25]: train_x.shape, test_x.shape

Out[25]: ((429, 24), (185, 24))

In [28]: from sklearn.tree import DecisionTreeClassifier
dt_model = DecisionTreeClassifier(max_depth=14)

In [29]: dt_model.fit(train_x, train_y)

Out[29]:

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=14)

In [30]: train_y_hat = dt_model.predict(train_x)
test_y_hat = dt_model.predict(test_x)

In [31]: print('-'*20,'Train','-'*20)
print(classification_report(train_y,train_y_hat))
print('-'*20,'Test','-'*20)
print(classification_report(test_y,test_y_hat))

```

----- Train -----
              precision    recall  f1-score   support

      0       0.99      1.00      1.00       127
      1       1.00      1.00      1.00       302

 accuracy      1.00
 macro avg     1.00
 weighted avg   1.00

----- Test -----
              precision    recall  f1-score   support

      0       0.56      0.54      0.55        65
      1       0.76      0.78      0.77       120

 accuracy      0.69
 macro avg     0.66
 weighted avg   0.69

```

In [32]: confusion_matrix(train_y, train_y_hat)

Out[32]: array([[127, 0],
[1, 301]], dtype=int64)

In []:

Practical No: 12

Aim: Program to Implement Random Forest

Output:

```
In [3]: import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
```

```
In [5]: train = pd.read_csv(r"D:\Siddharth_63\AIML\titanic - titanic.csv")
print(train.shape)

(891, 12)
```

```
In [6]: #checking for missing data
NAs = pd.concat([train.isnull().sum()], axis=1, keys=["Train"])
NAs[NAs.sum(axis=1) > 0]
```

```
Out[6]:
```

	Train
Age	177
Cabin	687
Embarked	2

```
In [7]: train.pop("Cabin")
train.pop("Name")
train.pop("Ticket")
```

```
Out[7]:
```

0	A/5	21171
1	PC	17599
2	STON/O2.	3101282
3		113803
4		373450
...		
886		211536
887		112053
888	W./C.	6607
889		111369
890		370376

Name: Ticket, Length: 891, dtype: object

```
In [8]: # Filling missing Age values with mean
train["Age"] = train["Age"].fillna(train["Age"].mean())
```

```
In [9]: # Filling missing Embarked values with most common values
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])
```

```
In [10]: train["Pclass"] = train["Pclass"].apply(str)
```

```
In [11]: # Getting Dummies from all other categories vars
for col in train.dtypes[train.dtypes == "object"].index:
    for dummy in train.pop(col):
        train = pd.concat([train, pd.get_dummies(for_dummy, prefix=col)], axis=1)
train.head()
```

```
Out[11]:
```

	PassengerId	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	1	0	22.0	1	0	7.2500	0	0	1	0	1	0	0	1
1	2	1	38.0	1	0	71.2833	1	0	0	1	0	1	0	0
2	3	1	26.0	0	0	7.9250	0	0	1	1	0	0	0	1
3	4	1	35.0	1	0	53.1000	1	0	0	1	0	0	0	1
4	5	0	35.0	0	0	8.0500	0	0	1	0	1	0	0	1

```
In [12]: labels = train.pop("Survived")
```

```
In [13]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(train, labels, test_size=0.25)
```

```
In [14]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100)
rf.fit(x_train, y_train)
```

```
Out[14]:
```

* RandomForestClassifier
 RandomForestClassifier()

```
In [15]: y_pred = rf.predict(x_test)
```


```
In [16]: from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

```
Out[16]: 0.798442760942761
```

```
In [17]: n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
```

```
In [18]: for estimator in n_estimators:
rf = RandomForestClassifier(n_estimators=estimator, n_jobs=-1)
rf.fit(x_train, y_train)
train_pred = rf.predict(x_train)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
train_results.append(roc_auc)
y_pred = rf.predict(x_test)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
test_results.append(roc_auc)
```

```
In [20]: from matplotlib.legend_handler import HandlerLine2D
line1, = plt.plot(n_estimators, train_results, "b", label="Train AUC")
line2, = plt.plot(n_estimators, test_results, "r", label="Test AUC")
plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
plt.ylabel("AUC score")
plt.xlabel("n_estimators")
plt.show()
```



```
In [19]: from sklearn.ensemble import RandomForestClassifier
rgf = RandomForestClassifier(n_estimators=200)
rgf.fit(x_train, y_train)
```

```
Out[19]:
RandomForestClassifier
RandomForestClassifier(n_estimators=200, n_jobs=-1)
```


```
In [20]: y_predict = rgf.predict(x_test)
from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_predict)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

```
Out[20]: 0.7986952861952863
```

```
In [21]: n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
```

```
In [22]: for estimator in n_estimators:
         rf = RandomForestClassifier(n_estimators=estimator, n_jobs=-1)
         rf.fit(x_train, y_train)
         train_pred = rf.predict(x_train)
         false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
         roc_auc = auc(false_positive_rate, true_positive_rate)
         train_results.append(roc_auc)
         y_pred = rf.predict(x_test)
         false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
         roc_auc = auc(false_positive_rate, true_positive_rate)
         test_results.append(roc_auc)
```

```
In [23]: from matplotlib.legend_handler import HandlerLine2D
         line1, = plt.plot(n_estimators, train_results, "b", label="Train AUC")
         line2, = plt.plot(n_estimators, test_results, "r", label="Test AUC")
         plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
         plt.ylabel("AUC score")
         plt.xlabel("n_estimators")
         plt.show()
```



Practical No: 13

Aim: Program to Implement AdaBoost

Output:

```
In [17]: from sklearn.ensemble import AdaBoostClassifier
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn import metrics
```

```
In [18]: #Load data
iris = datasets.load_iris()
X = iris.data
y = iris.target
```

```
In [19]: # Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
# 80% training and 20% test
```

```
In [20]: # Create adaboost classifier object
AdaModel = AdaBoostClassifier(n_estimators=100, learning_rate=1)

# Train Adaboost Classifier
model = AdaModel.fit(X_train, y_train)

# Predict the response for test dataset
y_pred = model.predict(X_test)
```

```
In [21]: # Model Accuracy, how often is the classifier correct?
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

Accuracy: 0.9333333333333333
```

Activate Win
Go to Settings to

```
In [22]: # Import Support Vector Classifier
from sklearn.svm import SVC
#Import scikit-learn metrics module for accuracy calculation
from sklearn import metrics
svc=SVC(probability=True, kernel='linear')
# Create adaboost classifier object
abc =AdaBoostClassifier(n_estimators=50, base_estimator=svc, learning_rate=1)
```

```
In [23]: # Train Adaboost Classifier
model = abc.fit(X_train, y_train)

#Predict the response for test dataset
y_pred = model.predict(X_test)
```

```
In [24]: # Model Accuracy, how often is the classifier correct?
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

Accuracy: 1.0
```


Practical No: 14

Aim: Program to Implement Gradient Boosting

Output:

```
In [1]: #Importing necessary packages
# Load libraries
from sklearn.ensemble import GradientBoostingRegressor
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.datasets import load_boston
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Load data - Reading Boston Data
boston = load_boston()
X = pd.DataFrame(boston.data, columns=boston.feature_names) #Independent columns

y = pd.Series(boston.target) #Dependent column - Median value of House
```

```
In [3]: #Viewing Data - predictors
X.head()
```

Out[3]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [4]: y[1:10] #response|
```

Out[4]:

1	21.6
2	34.7
3	33.4
4	36.2
5	28.7
6	22.9
7	27.1
8	16.5
9	18.9

dtype: float64

```
In [5]: # Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2) # 80% training and 20% test
```

```
In [6]: # Create gradientboost REGRESSOR object
gradientregressor = GradientBoostingRegressor(max_depth=2,n_estimators=3,learning_rate=1.0)
```

```
In [7]: # Train gradientboost REGRESSOR
model = gradientregressor.fit(X_train, y_train)

#Predict the response for test dataset
y_pred = model.predict(X_test)
```


```
In [8]: r2_score(y_pred,y_test)
```

Out[8]: 0.4839117761877929

```
In [9]: import matplotlib.pyplot as plt
%matplotlib inline

# Plot feature importance
feature_importance = model.feature_importances_

# make importances relative to max importance
feature_importance = 100.0 * (feature_importance / feature_importance.max())
sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + .5
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, boston.feature_names[sorted_idx])
plt.xlabel('Relative Importance')
plt.title('Variable Importance')
plt.show()
```



Activate
your license

```
In [10]: from sklearn.model_selection import GridSearchCV
LR = {'learning_rate':[0.15,0.1,0.10,0.05], 'n_estimators':[100,150,200,250]}

tuning = GridSearchCV(estimator=GradientBoostingRegressor(),
                      param_grid=LR, scoring='r2')
tuning.fit(X_train,y_train)
tuning.best_params_, tuning.best_score_
```

```
Out[10]: ({'learning_rate': 0.1, 'n_estimators': 150}, 0.8643406323195935)
```